Adversarial Generation and Collaborative Evolution of Safety-Critical Scenarios for Autonomous Vehicles Supplementary Material

Anonymous Author(s)
Affiliation
Address
email

Contents

Submitted to 39th Conference on Neural Information Processing Systems (NeurIPS 2025). Do not distribute.

A Prompting Strategy for Adversarial Scenario Generation

We design a two-stage prompting strategy that bridges natural language reasoning and executable simulation synthesis to enable automated generation of safety-critical scenarios. The first stage produces structured scenario descriptions involving a single adversarial agent, while the second translates these descriptions into partial Scenic code that defines spatial layout and behavior. This section details the prompt designs and their integration into the overall generation pipeline.

A.1 Instruction Prompt for Structured Scenario Description

In the first stage, we prompt the LLM to generate a structured description of a safety-critical scenario involving a single adversarial agent. The output includes a concise narrative and key semantic fields that define the agent's type, position, and unsafe behavior, forming the basis for later code generation. This step operates at the semantic level, where the model reasons plausible but dangerous traffic interactions grounded in real-world knowledge.

Prompt Overview. The instruction prompt is carefully engineered to guide the model's generative behavior while maintaining semantic fidelity and controllability. As shown in Tab. 1, the instruction prompt comprises five components: ❶ *Role Setting*, which defines the model's role as an expert scenario designer for AV safety testing; ❷ *Task Instructions*, which define the expected generative steps and constraints; ❸ *Knowledge Reference*, which specifies three external knowledge bases used to ground the reasoning in domain-relevant priors; ❹ *Output Format*, which enforces a structured, machine-readable response schema; and ❺ *Base Scenario Examples*, which provide in-context illustrations for the model to follow.

Design Rationale. The prompt is carefully structured to ensure semantic consistency and behavioral realism. Each generated scenario introduces a clear and interpretable threat aligned with known traffic violations and risk patterns by focusing on a single adversarial agent exhibiting unsafe behavior. Safety-related knowledge sources are essential: driving rules and licensing tests promote normative behavior baselines, while pre-crash scenarios provide realistic references for generating unsafe actions.

Template-Controlled Generation. During execution, the base scenario description is slotted into the `base_scenario` placeholder, allowing the LLM to contextualize its generation. The adversarial agent's type, position, and behavior are then extracted from the description and returned in a fixed schema (`AdvType`, `AdvPos`, `AdvBehavior`), ensuring seamless downstream parsing and programmatic conversion to Scenic code.

Few-shot Enhancement. To clarify the intent of each base scenario category and highlight its potential safety-critical variants, we provide a set of predefined examples as few-shot demonstrations. Rather than constraining generation, these examples serve as prompts that encourage the model to explore diverse and realistic adversarial situations beyond the given templates.

A.2 Scenic Prompt for Meta-Scenario Instantiation

In the second stage, we translate each structured scenario description into a partial Scenic program that specifies the adversarial agent's spatial configuration and unsafe behavior. This is achieved through a code-oriented prompt that guides the LLM in generating executable, simulator-compatible code. The resulting output defines a concrete meta-scenario that can be instantiated and evaluated in the CARLA environment.

Prompt Overview. As shown in Tab. 2, the Scenic prompt is systematically structured into five key sections, each serving a distinct purpose to guide the LLm toward generating high-quality, contextually accurate code. ❶ *Role Setting* establishes the model's identity as a domain expert in both Scenic programming and Carla-based simulation development. This section is crucial for contextual alignment, as it primes the model with relevant domain knowledge and sets expectations regarding the technical rigor required in its responses. ❷ *Task Instructions* explicitly describe the code generation objective, typically synthesizing Scenic programs based on high-level semantic descriptions of driving scenarios. This component ensures that the model understands the programming task at hand and the underlying semantics that the generated code should capture. ❸ *Goal* delineates the intended outcomes of the prompt, including the instantiation of autonomous agents, specification of their

Role Setting
You are now acting as an expert scenario designer for autonomous driving safety testing. ...
Task Instructions
Your task involves the following steps:

1. Understand the provided base scenario category. The provided base scenario is: [{base_scenario}]
2. Create one safety-critical scenario by introducing **ONE** adversarial traffic participant whose abnormal or rule-violating behavior leads to a collision or safety threat to the ego vehicle. *You may include other normal traffic participants to construct the scenario, but only ONE participant should behave adversarially.*
3. Extract and output the following elements from your generated scenario: AdvType: ...; AdvPos: ...; AdvBehavior: ...

Your generated description must be:

- Logically consistent with the given base scenario type
- Realistic and aligned with driving dynamics
- Focused on behavior that plausibly causes a collision or serious risk

Knowledge Reference
Use the following files as knowledge bases to ensure behavioral realism and rule awareness:

- driving_rules: ...
- dangerous_scenario: ...
- driving_test: ...

Incorporate these sources' reasoning patterns and behavior types when constructing your scenarios and describing adversarial actions.
Output Format
Respond strictly in the following format:
”’Text
Description: [Your scenario description in 24 concise and realistic sentences]
AdvType: [car / bicycle / motorcycle / pedestrian]
AdvPos: [e.g., "left front", "right rear", "behind", etc.]
AdvBehavior: [A short, clear description of the adversarial behavior]
”’
Do not modify field names. Your response must be machine-readable and follow this format exactly.
Base Scenario Examples
Here are two examples per category to help you understand how to construct safety-critical scenarios under different base scenarios. Focus on how the adversarial participant's position and behavior create danger:

1. Straight Obstacle
   Example 1: ...
   Example 2: ...
2. ...

Using the given base scenario, generate a safety-critical scenario and return the structured output.

Table 1: Instruction prompt used for structured safety-critical scenario description.

Role Setting
You are highly proficient in the Scenic programming language version 2.x and experienced in building Carla-compatible safety-critical driving scenarios based on real-world traffic behaviors. ...
Task Instructions
Your task is to understand the provided safety-critical scenario and translate it into a Scenic code snippet for Carla simulation.
The provided inputs are:

- {Description}: The complete description of a safety-critical scenario involving an ego vehicle and an adversarial participant.
- {AdvType}: The type of the adversarial traffic participant.
- {AdvPos}: The adversarial participant's relative position to the ego vehicle.
- {AdvBehavior}: The specific unsafe or abnormal behavior performed by the adversarial participant.

Your Goal
Based on the scenario description, generate a partial Scenic script that:

- Implements the adversarial behavior via a behavior block
- Defines the adversarial participant's position and properties

Note: The ego vehicle's spawn point (EgoSpawnPt) and trajectory (egoTrajectory) are already defined and do not need to be modified.
Scenic Code Template
Complete the Scenic snippet below by filling in the appropriate logic and values derived from Description, AdvType, AdvPos, and AdvBehavior:

```
'''scenic
Town = globalParameters.town
EgoSpawnPt = globalParameters.spawnPt
yaw = globalParameters.yaw
lanePts = globalParameters.lanePts
egoTrajectory = PolylineRegion(globalParameters.waypoints)

param map = localPath()
param carla_map = Town
model scenic.simulators.carla.model

EGO_MODEL = "vehicle.lincoln.mkz_2017"
ego = Car at EgoSpawnPt,
    with heading yaw,
    with blueprint EGO_MODEL

behavior AdvBehavior():
    TODO: Implement adversarial behavior based on [AdvBehavior]

TODO: Define adversarial agent position based on [AdvPos]
TODO: Choose an appropriate model for [AdvType] and complete the position and heading
AdvAgent = [AdvType] at xxx,
    with heading xxx,
    with behavior AdvBehavior()
'''
```

Do not modify field names. Your response must be machine-readable and follow this format exactly.
Output Format
Wrap scenic code in this format:

```
'''Scenic
[Full Scenic code here]
'''
```

Table 2: Prompt used to guide LLM-based generation of Scenic code for meta-scenario instantiation.

physical properties, and definition of their dynamic behaviors. By articulating concrete deliverables, this section serves as a guiding target that helps constrain the model's generation within meaningful and executable bounds. ❹ *Scenic Code Template* provides a structured scaffold of a Scenic program, often partially filled with placeholders or pre-defined elements. This template helps reduce ambiguity and encourages the model to focus on completing the code in a way that is syntactically valid and semantically appropriate for the described scenario. ❺ *Output Format*specifies strict requirements for the format and structure of the model's response, typically enforcing machine readability and syntactic correctness. This ensures that the generated Scenic code can be directly integrated into downstream simulation pipelines with minimal post-processing, thus promoting automation and robustness in scenario generation workflows. Together, these components form a comprehensive prompt design that instructs the model on what to generate and how and why, effectively bridging the gap between natural language scenario descriptions and executable simulation scripts.

Code Generation Design. The template is grounded in a standardized simulation setup, with map parameters, the ego vehicle's pose (`EgoSpawnPt`), and trajectory (`egoTrajectory`) predefined via `globalParameters`. The ego vehicle follows a fixed test route and is controlled by a black-box decision-making policy. In contrast, the adversarial agent must be fully specified in the generated code, including its spatial position and behavior logic. The agent is defined by its type (`AdvType`), positioned relative to the ego vehicle according to the semantic field `AdvPos`, and mapped to global coordinates using parametric anchor points (*e.g.*, `OPT_X_DIST`, `OPT_Y_DIST`) to ensure physical plausibility. Its behavior is implemented as a standalone `behavior` block derived from the natural language field `AdvBehavior`. This design allows clear separation between agent instantiation and control, enabling Scenic's reactive modeling interface to express temporally coordinated and semantically grounded unsafe actions.

B Scenario Instantiation and Adaptive Parameter Optimization

To illustrate the generation and selection process of adversarial scenarios in our framework, we present a case based on Scenario 1 and the route numbered 9. This case is randomly chosen for demonstration purposes. The steps described below, which include Scenic-based instantiation, iterative parameter refinement, and simulation-based scene selection, are applied consistently across all scenarios and route combinations used for evaluation.

B.1 Generated Meta-Scenario represented by Scenic Code

Given a base scenario of type `Straight Obstacle`, where the ego vehicle travels along a straight urban road, the LLM is prompted to generate a safety-critical meta-scenario by introducing a single adversarial agent. The model outputs a structured description in the form

$$\langle \Phi_{\mathrm{c}}, \Phi_{\mathrm{p}}, \Phi_{\mathrm{b}}, \Phi_{\mathrm{R}}, \Phi_{\mathrm{L}} \rangle, \tag{B.1}$$

where each Φ_* represents a natural language element describing the scenario's configuration. Specifically, Φ_{c}, Φ_{p}, and Φ_{b} denote the adversarial agent's type, position, and behavior. Φ_{R} refers to the road structure relevant to the ego vehicle, and Φ_{L} captures any applicable traffic signal state. For this case, the generated values are shown in Tab. 3.

Based on this structured output, the language model generates an executable Scenic script that instantiates the meta-scenario for simulation, as shown in Tab. 4. This script encodes a semantically grounded meta-scenario by defining the ego vehicle and the adversarial agent with detailed control over spatial configuration, behaviors, and parametric uncertainty. The ego vehicle is instantiated at a predefined spawn point (`EgoSpawnPt`) with a fixed heading and a prescribed trajectory (`egoTrajectory`). These predefined values are constructed based on the NHTSA Pre-Crash Typology Report. The adversarial agent is instantiated as a pedestrian in this example, and is placed relative to a blocking truck (`Blocker`) located at the right front of the ego vehicle, using a directional offset computed from road geometry.

The blocking truck conceals the pedestrian and introduces realistic occlusion, increasing the risk of delayed perception or response by the ego vehicle. The pedestrian is placed behind the blocker, facing perpendicular to the direction of travel (rotated by $+90^\circ$ relative to the road heading), and its motion is governed by a custom-defined `behavior` function, constructed according to Scenic's reactive

Φ_{base}: Straight Obstacle
Φ_{meta}: The ego vehicle travels straight when a pedestrian suddenly steps out from behind a parked car on the right front, darts into the roadway, and abruptly stops mid-path directly in the ego's trajectory.

Φ_{c} (AdvType): pedestrian
Φ_{p} (AdvPos): right front
Φ_{b} (AdvBehavior): Emerges from concealment, crosses unexpectedly while freezing in the path
Φ_{R} (road layout): straight road
Φ_{L} (traffic signal): None

Table 3: Example output of the instruction prompt for structured scenario description.

behavior model. In Scenic, a `behavior` is a composition of conditionally triggered actions or nested sub-behaviors, where actions are executed via `take` and behaviors via `do`. The `AdvBehavior` defined here first invokes the built-in `CrossingBehavior(...)`, which is a synchronization behavior that adjusts the pedestrian's speed to ensure intersection with the ego vehicle's trajectory at a critical moment. This crossing logic dynamically regulates the agent's velocity based on distance and expected arrival timing. After executing the crossing, the agent transitions into a passive state by continuously taking a zero-speed walking action (`SetWalkingSpeedAction(0)`), simulating a sudden freeze in the ego vehicle's path.

All behavioral and spatial parameters involved in the meta-scenario, such as the pedestrian's walking speed, lateral offset from the blocker, and stopping distance relative to the ego vehicle's trajectory, are defined using `Range(...)` expressions in Scenic. Each parameter represents a probability distribution rather than a fixed value, which allows the simulation to sample diverse and realistic scene variations. This probabilistic modeling approach supports systematic exploration of the scenario space.

B.2 Adaptive Parameter Optimization

After defining the scenario in Scenic, we introduce adaptive optimization to refine key behavioral and spatial parameters, such as walking speed, emergence offset, and stopping distance. These parameters are initially defined as probabilistic `Range(...)` expressions, enabling diverse sampling of simulation scenario instances.

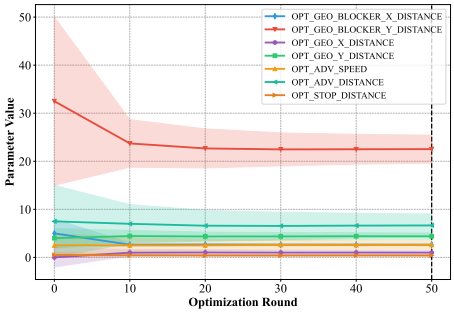

Figure 1: Refinement of parameter distributions across optimization rounds.

To ensure that the sampled configurations effectively induce safety-critical outcomes, we iteratively adjust the parameter distributions based on feedback from simulation results. Each meta-scenario is simulated fifty times in the CARLA environment, where parameter values are drawn from their initial predefined ranges. The process is organized into five intervals of ten runs each. After every interval, we analyze the simulation outcomes and identify those runs in which a collision involving the ego vehicle occurs.

We collect the values corresponding to these collision-inducing runs for each parameter and estimate the empirical mean μ and standard deviation σ. In the next optimization round, the sampling range is updated to $[\mu - \sigma, \mu + \sigma]$, concentrating the search in regions associated with higher risk while maintaining behavioral realism.

This iterative refinement process gradually narrows the parameter space toward semantically consistent and behaviorally disruptive configurations. After all fifty runs, we rank the resulting scene instances based on collision severity metrics, such as time to collision, trajectory deviation, and performance degradation of the autonomous driving system. The top K scenes with the most substantial

```
Town = globalParameters.town
EgoSpawnPt = globalParameters.spawnPt
yaw = globalParameters.yaw
egoTrajectory = PolylineRegion(globalParameters.waypoints)
param map = localPath(f'../maps/Town.xodr')
param carla_map = Town
model scenic.simulators.carla.model
EGO_MODEL = "vehicle.lincoln.mkz_2017"
ADV_MODEL = "walker.pedestrian.0009"
BLOCKER_MODEL = "vehicle.volkswagen.t2"

ego = Car at EgoSpawnPt,
    with heading yaw,
    with regionContainedIn None,
    with blueprint EGO_MODEL

param OPT_BLOCKER_X = Range(2, 8)
param OPT_BLOCKER_Y = Range(15, 50)
param OPT_GEO_X = Range(-2, 2)
param OPT_GEO_Y = Range(2, 6)
param OPT_ADV_SPEED = Range(0, 5)
param OPT_ADV_DISTANCE = Range(0, 15)
param OPT_STOP_DISTANCE = Range(0, 1)

behavior AdvBehavior():
        do CrossingBehavior(ego, globalParameters.OPT_ADV_SPEED, globalParameters.OPT_ADV_DISTANCE) until (distance from self to egoTrajectory) < globalParameters.OPT_STOP_DISTANCE
    while True:
        take SetWalkingSpeedAction(0)

IntSpawnPt = OrientedPoint following roadDirection from EgoSpawnPt for globalParameters.OPT_BLOCKER_Y

Blocker = Car right of IntSpawnPt by globalParameters.OPT_BLOCKER_X,
    with heading IntSpawnPt.heading,
    with regionContainedIn None,
    with blueprint BLOCKER_MODEL

SHIFT = globalParameters.OPT_GEO_X @ globalParameters.OPT_GEO_Y
AdvAgent = Pedestrian at Blocker offset along IntSpawnPt.heading by SHIFT,
    with heading IntSpawnPt.heading + 90 deg,
    with regionContainedIn None,
    with behavior AdvBehavior(),
    with blueprint ADV_MODEL

require (distance from AdvAgent to intersection) > 10
```

Table 4: Scenic definition of a meta-scenario for simulation in CARLA.

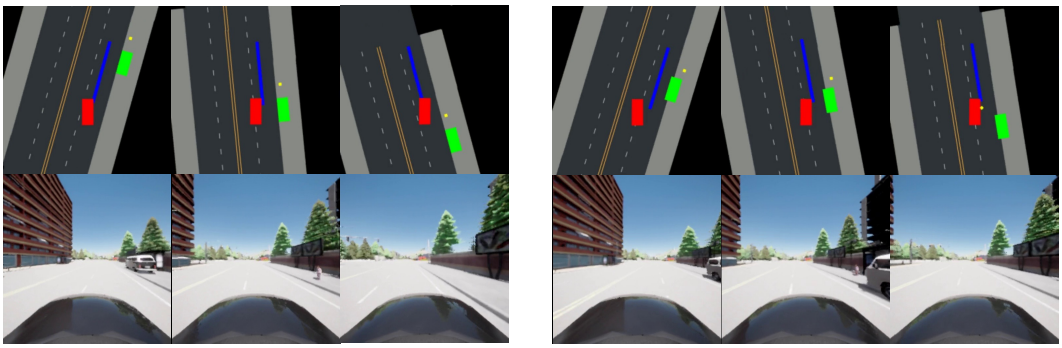

(a) Discarded instance. The adversarial pedestrian reacts too late, initiating crossing only after the ego vehicle has passed, resulting in a non-critical scenario.

(b) Retained instance. The pedestrian emerges at a critical moment, producing a high-risk interaction that leads to a collision or forced maneuver.

Figure 2: Qualitative comparison of discarded and retained scene instances derived from the same Scenic specification. Parameter optimization shifts the adversarial behavior from low to high impact while preserving semantic structure.

adversarial impact are selected for final evaluation. Our experiments use $K = 2$ per route, resulting in a curated suite of 800 adversarial scenarios.

This approach supports the generation of high-risk yet physically and semantically plausible scenarios. Compared to direct optimization-based attacks, it maintains the diversity and interpretability of the scene space while effectively exposing weaknesses in autonomous vehicle policies.

We present qualitative comparisons between representative scenario instances selected and discarded during optimization to support the quantitative evidence of parameter refinement. These examples visually demonstrate the ability of our iterative strategy to isolate high-risk configurations from a broad range of plausible parameter combinations.

To further illustrate the behavioral divergence caused by parameter refinement, we provide a visual comparison in Fig. 2. The two subfigures represent scenario instances instantiated from the same Scenic meta-scenario template, but sampled at different stages of the optimization process.

In Fig. 2a, the adversarial pedestrian fails to induce a safety-critical interaction. Due to the unrefined setting of the trigger distance, the pedestrian begins crossing only after the ego vehicle has passed the occluded region. This results in a benign interaction with no evasive behavior required.

By contrast, Fig. 2b shows a retained instance where the parameters have been refined to achieve tighter spatial and temporal alignment. Here, the pedestrian emerges from concealment when the ego vehicle approaches the crossing path, creating a high-risk situation with minimal reaction time. This scenario is successfully flagged as adversarial and selected for final evaluation.

These qualitative differences highlight the impact of outcome-guided parameter optimization in shifting agent behavior from neutral to adversarial, even when the underlying scenario semantics remain unchanged. This demonstrates that parameter tuning alone can lead to substantially different risk profiles in simulation.

B.3 Scenic-to-CARLA Integration

To operationalize meta-scenarios generated in Scenic, we develop a simulation integration pipeline that bridges high-level semantic descriptions with executable behavior in the CARLA environment. Scenic is an intermediate, declarative specification language that enables structured scene definitions with precise spatial layouts, probabilistic variation, and reactive agent behavior.

Each Scenic script is compiled into a Python execution plan via the Scenic runtime, which interacts directly with CARLA's Python API. This process resolves object placement, behavior execution, and physical constraints at runtime. The ego vehicle is instantiated using a fixed blueprint and global spawn point, following a predefined route encoded as a set of waypoints. Its closed-loop control

is governed by a black-box driving policy, which consumes fused sensor observations and outputs low-level control commands.

Dynamic agents, including adversarial participants and background vehicles, are instantiated according to their Scenic-defined geometry and temporal behavior logic. These agents execute behaviors defined through conditionally triggered actions and sub-behaviors, translated into low-level simulator commands via Scenic-CARLA bindings. Crucially, the parameterized behaviors are resolved through sampled values from `Range(...)` expressions, yielding diverse scene instantiations.

Spatial validity is ensured using road topology information extracted from the underlying map (*e.g.*, OpenDRIVE format). For example, statements like 'right of intersection by offset' are grounded via directional offsets aligned with local lane orientation. All agent placements and motion constraints adhere to physical feasibility and collision-free initialization.

During simulation, each run corresponds to a single sampled realization of the scenario. The simulator operates in synchronous mode, enabling frame-accurate coordination between the ego vehicle and surrounding agents. Per-frame logging of collisions, trajectory deviation, rule violations, and behavioral states supports downstream evaluation and adaptive optimization.

This integration module is essential to our SCENGE pipeline. It ensures that semantically grounded, LLM-generated scenarios can be faithfully instantiated and executed in a closed-loop environment, supporting scalable testing of autonomous vehicle behavior under realistic and safety-critical conditions.

C Attention Matrix and Temporal Modeling

We introduce a temporal attention mechanism that quantifies the relevance between dynamic agents over time to identify influential background agents that may contribute to safety-critical interactions. The goal is to move beyond static heuristics such as proximity or time-to-collision, and instead leverage temporal cross-agent patterns to capture more nuanced, causally grounded dependencies.

We define an attention matrix $\boldsymbol{M}_{\text{att}} \in \mathbb{R}^{(2T)\times(NT)}$, where T denotes the number of time frames, and N is the number of background vehicles. Each row corresponds to a specific timestep of either the ego vehicle or the adversarial agent (*i.e.*, T rows for each), while each column corresponds to a particular frame of a background agent. The entry $\mathbf{M}_{\text{att}}(t, j)$ encodes the relevance of background agent $a_{\lfloor j/T \rfloor}$ at frame $j \bmod T$ to the querying agent at frame t.

The structure is illustrated in Eq. (C.1), where blue rows denote frames of the ego vehicle, green columns indicate background agent $\boldsymbol{\tau}_1$, and the red region highlights the cumulative relevance of background agent $\boldsymbol{\tau}_2$ across frame *w.r.t* the ego vehicle.

$$
\begin{array}{c}
\\
\tau_{\text{ego}}(0) \\
\tau_{\text{ego}}(1) \\
\vdots \\
\tau_{\text{ego}}(T-1) \\
\tau_{\text{adv}}(0) \\
\vdots
\end{array}
\begin{array}{c}
\overbrace{\phantom{m_{1,1}\cdots m_{1,T}}}^{\tau_1(0)\quad \tau_1(1)\quad \tau_1(T-1)} \quad \overbrace{\phantom{m_{1,T+1}\cdots m_{1,2T}}}^{\boldsymbol{\tau}_2} \quad \tau_N(T-1) \\
\begin{pmatrix}
m_{1,1} & \cdots & m_{1,T} & m_{1,T+1} & \cdots & m_{1,2T} & \cdots & m_{1,NT} \\
 & & & & & & & \vdots \\
m_{T,1} & \cdots & & m_{T,T+1} & \cdots & m_{T,2T} & \cdots & m_{T,NT} \\
m_{T+1,1} & \cdots & & & & & & m_{T+1,NT} \\
\vdots & & & & & & & \vdots \\
m_{2T,1} & \cdots & & & & & & m_{2T,NT}
\end{pmatrix}
\end{array}
\tag{C.1}
$$

To ensure temporal consistency, we apply two auxiliary components:

- **Temporal Masking:** A binary mask $M_{\text{mask}}(t, j)$ sets attention to zero when $j \bmod T > t$, disallowing access to future frames.

$$
\boldsymbol{M}_{\text{mask}}(t, j) = \begin{cases} 0, & \text{if } j \bmod T \leq t \\ -\infty, & \text{otherwise} \end{cases} \tag{C.2}
$$

- **Temporal Decay:** A decay term down-weights distant past frames, where $\gamma \in (0, 1]$ controls the decay rate.

$$\boldsymbol{M}_{\text{decay}}(t, j) = \gamma^{t-(j \bmod T)}, \quad \gamma \in (0, 1] \tag{C.3}$$

After computing $\mathbf{M}_{\text{att}}$, we aggregate across rows corresponding to the ego vehicle to obtain a cumulative relevance score for each background vehicle. Vehicles with the highest cumulative scores are selected as collaborative perturbation candidates in the scenario evolution stage. This data-driven approach allows SCENGE to reason about multi-agent interactions and prioritize perturbations that are more likely to trigger unsafe outcomes in coordination with the primary adversary.

In contrast to purely proximity-based heuristics, our attention-based mechanism captures both temporal dynamics and asymmetric influence, allowing for more targeted and effective scenario manipulation.

D Loss Function Design and Behavioral Rationale

While the main paper defines the objective used for optimizing the trajectories of background collaborators, it does not elaborate on the design motivations underlying its components. We provide a detailed rationale for each term, highlighting its behavioral implications and geometric intent within the context of autonomous vehicle planning under uncertainty.

The first component encourages spatial proximity between the perturbed agent and the ego vehicle. Crucially, the goal is not merely to trigger collisions, but to induce planning instability by creating a perception of physical constraint. When an agent is placed within the ego's immediate operational envelope, especially in adjacent lanes or merging zones, the AV policy must resolve complex trade-offs between maintaining nominal progress and avoiding potential side conflicts. This is especially disruptive under partial observability, where predictions about intent and future occupancy are inherently uncertain. Thus, spatial closeness becomes a proxy for generating decision-theoretic ambiguity.

The second component targets occlusion-aware perturbation. By minimizing the orthogonal distance between the collaborator–ego and adversary–ego vectors, the optimization favors geometric alignments that maximize visual or LiDar occlusion. Such configurations simulate real-world cases where pedestrians emerge from behind parked vehicles or a vehicle masks another agent at intersections. Importantly, these are not direct collisions but structured perceptual traps that stress-test the perception-planning interface. Empirically, we observe that this term increases failure rates among models with strong reactive behavior but weak long-horizon anticipation.

The third term serves a dual role: it ensures that the trajectory perturbations remain dynamically feasible and that the resulting behaviors are behaviorally and statistically plausible to an AV model trained on naturalistic data. Without this regularization, agents may exhibit sudden velocity jumps or oscillatory paths that are physically implausible and trivially detected as out-of-distribution anomalies. From a systems perspective, ensuring smooth, non-suspicious motion allows failures to manifest from genuine planning or perception limitations, rather than from artificial data artifacts.

Beyond individual terms, the joint objective is designed to produce collaborators that function as implicit adversaries: not by violating rules directly, but by shaping the traffic context in ways that compound the risk induced by a designated primary adversary. This approach aligns with our broader goal of studying emergent multi-agent risks, where individually benign behaviors, when strategically composed, give rise to system-level safety violations.

By balancing proximity-induced pressure, occlusion-based ambiguity, and behavioral realism, our loss design facilitates the generation of diverse yet impactful scenario variants. These are particularly effective at revealing failure modes in vision-language and reinforcement learning–based autonomous agents that may otherwise appear robust under conventional single-agent test settings.

E Composite Metric Definitions

E.1 Metrics for VLM-Based Models

For instruction-following driving systems such as LMDrive, we adopt three canonical evaluation metrics as proposed in its original benchmark:

- **Route Completion (RC):** Measures the percentage of the planned route that the ego vehicle successfully traverses before a terminal failure (*e.g.*, collision or getting stuck). Higher RC indicates better progress and robustness.
- **Infraction Score (IS):** Penalizes unsafe behaviors during driving, including collisions, lane invasions, red-light violations, and traffic infractions. A lower IS reflects safer driving.
- **Driving Score (DS):** A composite metric that balances RC and IS, providing a joint measure of both task completion and safety. It is calculated using a predefined weighted formula from the LMDrive benchmark.

While these metrics offer a useful first-order evaluation of driving policy robustness, they are primarily designed for nominal driving settings. They may not fully reflect performance under targeted adversarial scenarios, where nuanced differences in perception, reaction, and trajectory adaptation are critical. To better capture these aspects, we introduce three additional composite metrics tailored to adversarial evaluation.

E.2 Three Composite Metrics

To support fine-grained evaluation of autonomous vehicle behavior under adversarial conditions, we define three composite metrics: *Safety Score*, *Task Score*, and *Comfort Score*. These scores are designed to reflect different aspects of policy robustness, including risk avoidance, task completion, and motion quality. Before defining the score formulas, we describe the component indicators used in their computation. These indicators are grouped into three behavior dimensions.

- **Safety-related indicators:**
  - **CR (Collision Rate)**: frequency of collisions involving the ego vehicle (↑)
  - **RR (Red Light Running)**: frequency of red light violations (↑)
  - **SS (Stop Sign Violation)**: frequency of failure to stop at stop signs (↑)
  - **OR (Off-Road Deviation)**: average distance driven outside the road boundary (↑)
- **Functionality-related indicators:**
  - **RF (Route Following Instability)**: deviation from the instructed navigation path (↓)
  - **Comp (Completion Ratio)**: percentage of the planned route completed (↓)
  - **TS (Time to Completion)**: time required to finish the route (↑)
- **Etiquette-related indicators:**
  - **ACC (Acceleration)**: average longitudinal acceleration (↑)
  - **YV (Yaw Velocity)**: average angular velocity while turning (↑)
  - **LI (Lane Invasion)**: frequency of unintended lane departures (↑)

Higher values marked with ↑ indicate degraded behavior, while lower values marked with ↓ indicate improved behavior. All metrics are normalized to the range $[0, 1]$ using the following formulas. Higher scores indicate better performance by the autonomous vehicle when responding to the given scenario. Conversely, lower scores indicate greater behavioral degradation and are interpreted as a proxy for stronger adversarial effect.

Safety Score evaluates the vehicle's ability to avoid direct collisions and latent safety risks. It is defined as

$$\mathrm{OS}_{\mathrm{Safety}} = \frac{5}{8}(1 - \mathrm{CR}) + \frac{1}{8}(3 - \mathrm{RR} - \mathrm{SS} - \mathrm{OR}), \tag{E.1}$$

A higher Safety Score reflects lower collision frequency and better compliance with basic traffic safety constraints.

Task Score measures how well the agent follows the navigation plan and completes its assigned task. It is defined as

$$\mathrm{OS}_{\mathrm{Task}} = \frac{1}{3}(\mathrm{RF} + \mathrm{Comp} + 1 - \mathrm{TS}), \tag{E.2}$$

A higher Task Score corresponds to better route adherence, higher completion percentage, and less time required to finish.

Comfort Score captures motion smoothness and stability in trajectory execution. It is defined as

$$\mathrm{OS}_{\mathrm{Comfort}} = \frac{1}{3}(3 - \mathrm{ACC} - \mathrm{YV} - \mathrm{LI}). \tag{E.3}$$

This score penalizes aggressive maneuvers, oscillatory motion, and unintended lane departures. A higher score indicates smoother and more stable driving. We report all three scores for each adversarial scenario to enable a multidimensional assessment of policy robustness. These metrics help characterize the behavioral degradation induced by semantically grounded adversarial perturbations and support consistent comparison across different scenarios and policies.

F Visual Illustration of Safety-Critical Scenarios

We provide representative snapshots from all eight base scenarios to visually demonstrate the adversarial impact and diversity of scenarios generated by SCENGE.

As shown in Fig. 3, the adversarial agent is a pedestrian concealed behind a parked truck on the right side of the road, exploiting a typical occlusion scenario to remain undetected by the ego vehicle until the last moment. As the ego vehicle proceeds along a straight path at a constant speed, the pedestrian suddenly emerges into its trajectory, triggering a safety-critical interaction with extremely limited reaction time. The pedestrian's emergence is precisely timed to occur when the ego vehicle has already committed to its motion, minimizing the opportunity for deceleration or evasive steering and thereby inducing a high-probability collision. This primary adversarial behavior is further exacerbated by the spatial configuration of background traffic: a vehicle traveling directly ahead of the ego car partially blocks the field of view and intensifies the occlusion caused by the parked truck. In contrast, another vehicle on the left side of the ego vehicle occupies the adjacent lane, preventing any potential escape maneuver through lane changing. Although these background agents do not behave adversarially themselves, their positions and trajectories contribute to a constrained and ambiguous environment that limits the ego vehicle's response options. The resulting collision highlights the vulnerability of autonomous systems to multi-agent interactions.

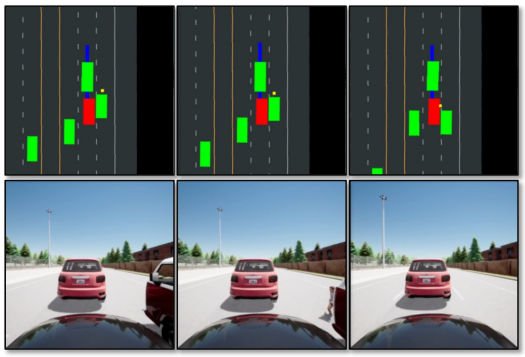

Figure 3: **Base Scenario 1:** Straight Obstacle.

This adversarial agent is a motorcycle as shown in Fig. 4. While the ego vehicle prepares to make an unprotected left turn at an unsignalized intersection, the motorcycle suddenly accelerates into the intersection from the opposite direction, exploiting the typical hesitation period during which the ego vehicle must infer the intentions and speed of oncoming traffic. The ego vehicle initiates the turn based on a perceived gap in traffic, but the motorcycle's rapid and unexpected approach collapses that gap before the maneuver can be completed safely. A background vehicle located at the right-front crossroad yields to the motorcycle, exhibiting behavior that would support safe intersection navigation under ordinary circumstances. However, in this case, yielding places the background vehicle in a position that

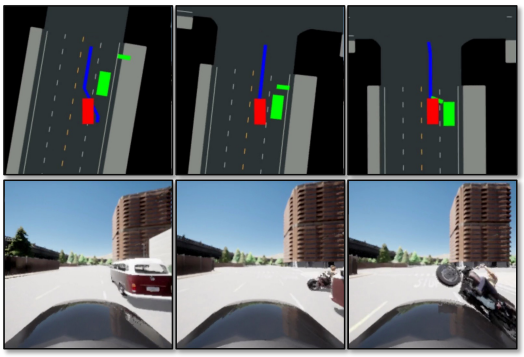

Figure 4: **Base Scenario 2:** Turning Obstacle.

inadvertently blocks the ego vehicle's view of the motorcycle's approach, creating a partial occlusion that critically delays detection. As a result, the ego vehicle's planning system does not register the threat until the motorcycle has entered the collision zone, leaving insufficient time to brake or adjust course. The collision that ensues demonstrates how even cooperative or rule-following actions by surrounding agents can, when spatially misaligned, amplify the impact of adversarial behavior. This scenario underscores the challenge of ensuring robust situational awareness in dense traffic environments and highlights the compound risk introduced by dynamic occlusions in multi-agent interactions.

As shown in Fig. 5, the ego vehicle attempts to perform a lane change when it is suddenly intercepted by an adversarial vehicle approaching from the right that initiates a nearly simultaneous maneuver into the same target lane, leading to a direct and immediate conflict over shared space. The adversarial vehicle's trajectory is aligned to converge with the ego vehicle's path at a critical point during the lane change, leaving almost no margin for temporal or spatial adjustment. A background vehicle following closely behind the ego car makes the situation more hazardous, which significantly restricts its ability to decelerate or abort the lane change by returning to its original position. This trailing vehicle does not exhibit aggressive behavior, but its proximity effectively removes any viable backward escape strategy for the ego vehicle. The resulting spatial entrapment leaves the ego system with no safe course of action, culminating in a side-impact collision that occurs precisely because of the compounded constraints imposed by multiple agents. This scenario illustrates the intricate risk landscape during routine driving behaviors such as lane changes, particularly when adversarial timing coincides with benign but spatially influential behaviors from surrounding vehicles. It highlights how even non-malicious background traffic can unintentionally support the success of adversarial actions by closing off critical degrees of freedom in the ego vehicle's decision space.

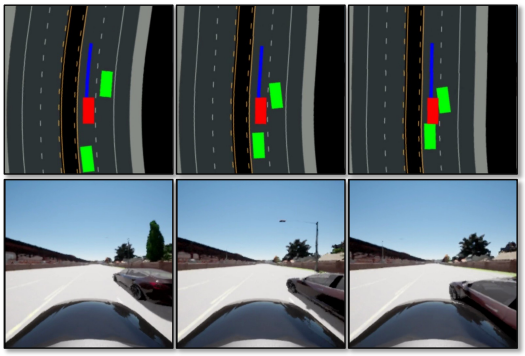

Figure 5: **Base Scenario 3:** Lane Changing.

As shown in Fig. 6, the ego vehicle attempts to overtake a slow-moving truck ahead by steering into the opposite lane to complete the pass within what is perceived to be a sufficient spatial and temporal gap. This maneuver is initiated based on the assumption that the oncoming lane is clear, a judgment made under conditions of limited visibility caused by the presence and size of the obstructing truck. However, as the ego vehicle commits to the overtaking action, an adversarial vehicle traveling at high speed from the opposite direction rapidly enters the scene, dramatically reducing the available time for evasive decision-making. The oncoming vehicle had not been within the ego vehicle's initial field of view due to topographical occlusion or the limited field of sensor perception, making its emergence effectively unanticipated. As the two vehicles close in distance rapidly, the ego vehicle is left with insufficient space to complete the overtaking maneuver and cannot safely retreat into its original lane because the slow-moving truck continues to occupy it. The result is a head-on collision occurring in a context where neither complete visibility nor real-time reactivity can compensate for the adversarially orchestrated timing of the opposing vehicle. This scenario underscores the acute risks associated with overtaking in partially observable environments. It illustrates how adversarial agents can precisely exploit the narrow margins within which such maneuvers operate to provoke failure. It reveals the vulnerability of autonomous systems when spatial assumptions are violated by rapid environmental changes that exceed the ego vehicle's planning horizon.

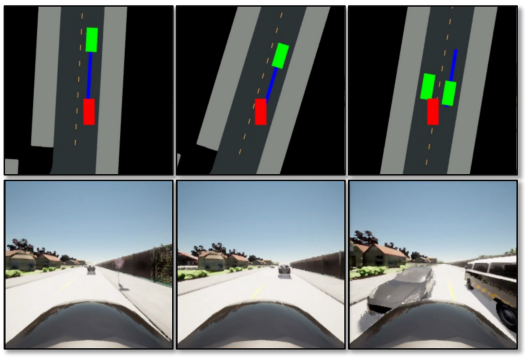

Figure 6: **Base Scenario 4:** Vehicle Passing.

As illustrated in Fig. 7, the ego vehicle proceeds straight through a signalized intersection with a green light, following traffic rules and anticipating a clear right-of-way. However, an adversarial vehicle approaches from the left side of the intersection. It violates the red light by entering the crossing at a significant speed, creating an immediate and severe threat to the ego vehicle's trajectory. The danger is compounded by a background vehicle traveling directly ahead of the ego car in the same lane. Although not adversarial, this vehicle limits the ego vehicle's field of vision. It restricts its ability to observe lateral traffic activity, particularly from the left side where the adversarial vehicle originates. As a result, the ego vehicle experiences a delay in detecting the oncoming threat and has minimal space and time to initiate any effective evasive maneuver. The forward vehicle also constrains the ego vehicle's longitudinal flexibility, limiting its capacity to accelerate or decelerate freely in response to sudden changes in the intersection environment.

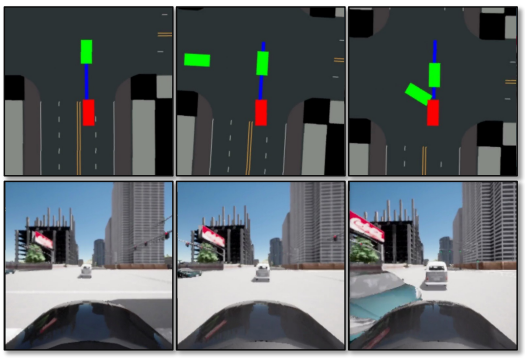

Figure 7: **Base Scenario 5:** Red-light Running.

In Fig. 8, the ego vehicle initiates an unprotected left turn at an intersection. As the vehicle progresses through the intersection, an adversarial vehicle from the opposite direction suddenly accelerates into the scene, targeting the point of conflict with precise timing that leaves the ego vehicle with minimal time to react. The adversarial vehicle's late-stage acceleration is particularly disruptive, where the ego vehicle must assess dynamic threats in real time while already partially exposed in the intersection. The situation is further complicated by a background vehicle traveling directly ahead of the ego car, whose presence partially obstructs the ego vehicle's view of the oncoming lane and conceals the initial approach of the adversarial vehicle. This occlusion prevents the ego system from recognizing the oncoming threat earlier. Because the ego vehicle has already entered the turning arc, its maneuverability is constrained, and its options are reduced to abrupt braking or incomplete evasive action, which is neither sufficient in the available time window.

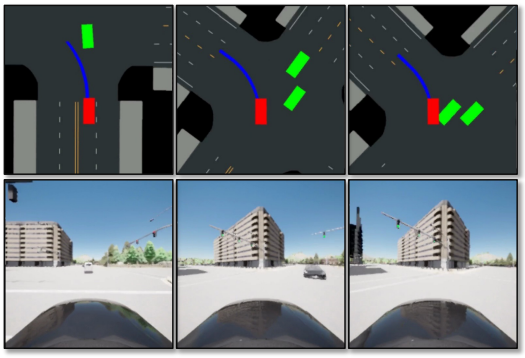

Figure 8: **Base Scenario 6:** Unprotected Left-turn.

As shown in Fig. 9, the ego vehicle performs a right turn at an intersection following a planned low-speed trajectory appropriate for a standard maneuver in structured traffic environments. While the turn is underway and the vehicle is partially through the arc, an adversarial vehicle rapidly approaches from the left with high velocity and an unstable trajectory that causes it to lose control and deviate from its expected path, cutting diagonally across the intersection in an abrupt and unanticipated manner. The ego vehicle, having already committed to the turning motion, operates within a narrow range of feasible control commands due to the low-speed dynamics and constrained steering angle typical of turning behavior in confined urban spaces. As a result, it lacks the capacity to execute a timely

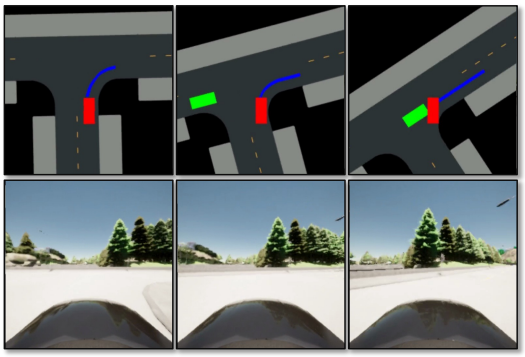

Figure 9: **Base Scenario 7:** Right-turn.

evasive action or halt its motion without causing further instability. The sudden intrusion of the adversarial vehicle from the lateral field, combined with the absence of sufficient temporal or spatial margin for re-planning, leads to a collision that cannot be mitigated through conventional reactive strategies. This scenario highlights the susceptibility of turning maneuvers to lateral disruptions from directions not directly aligned with the ego vehicle's forward-facing sensors or primary field of attention. It demonstrates how adversarial agents that exhibit unpredictable speed profiles and non-normative trajectories can exploit moments of limited agility and reduced observability during cornering actions, presenting a significant challenge to the robustness of trajectory prediction and contingency planning mechanisms within autonomous systems.

In the scenario illustrated by Fig. 10, the ego vehicle engages in a crossing negotiation at a multi-way intersection, proceeding cautiously while monitoring the behaviors of surrounding traffic participants to determine a safe moment to continue. During this process, an adversarial vehicle approaches from the opposite direction at a speed and trajectory, creating an imminent threat not directly to the ego vehicle but to another background vehicle traveling laterally across the intersection. The adversarial vehicle initiates a high-speed collision with the background vehicle, striking it with enough force to alter its heading and momentum. As a result of this initial impact, the background vehicle's front end is deflected off its original path and redirected into the lane occupied by the ego vehicle, creating an unexpected and unavoidable secondary collision. This sequence of events introduces a complex causal chain in which the ego vehicle becomes the victim of an indirect threat, which originates from the interaction between two other agents. The scenario exemplifies how adversarial behavior can manifest through direct engagement with the ego vehicle and the manipulation of nearby agents in a shared environment to create emergent hazards. The timing, positioning, and velocities of all involved vehicles contribute to the cascading nature of the incident, making it extremely difficult for the ego vehicle's planning and prediction modules to anticipate the eventual outcome based solely on initial observations. This case highlights the limitations of conventional threat models that focus primarily on direct interactions and emphasizes the necessity for autonomous systems to reason about second-order consequences of observed behaviors, particularly in densely populated or complex intersection settings where multi-agent dynamics can lead to rapidly evolving risk profiles.

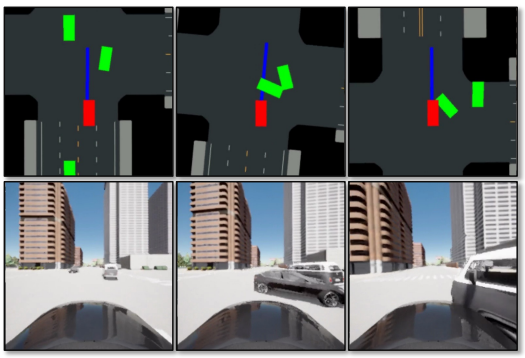

Figure 10: **Base Scenario 8:** Crossing Negotiation.